

Introduction to machine-learning for neuroimaging

May 2024

By

Qing Wang (Vincent)

王庆

<https://github.com/Vincent-wq>



国家精神疾病医学中心
NATIONAL CENTER FOR MENTAL DISORDERS
上海市精神卫生中心
SHANGHAI MENTAL HEALTH CENTER



Contents

- Part 1: Machine Learning Basics
 - Describe basics of the “learning” process
 - Explain model design choices and performance trade-offs
 - Model selection and validation frameworks
- Part 2: Machine Learning for Neuroimaging
 - Python lib family for neuroimaging studies
 - *Coding example: Autism fMRI classification*

What is machine learning?

What is machine learning?

Quick Internet search: *It is a branch of artificial intelligence (AI) and computer science that focuses on the using data and algorithms to enable AI to imitate the way that humans learn.*

What is machine learning?

Quick Internet search: *It is a branch of artificial intelligence (AI) and computer science that focuses on the using data and algorithms to enable AI to imitate the way that humans learn*

Philosophy professor: *You do not control what you make but you control the process of its creation.*


What is machine learning?

Quick Internet search: *It is a branch of artificial intelligence (AI) and computer science that focuses on the using data and algorithms to enable AI to imitate the way that humans learn.*

Philosophy professor: *You do not control what you make but you control the process of its creation.*

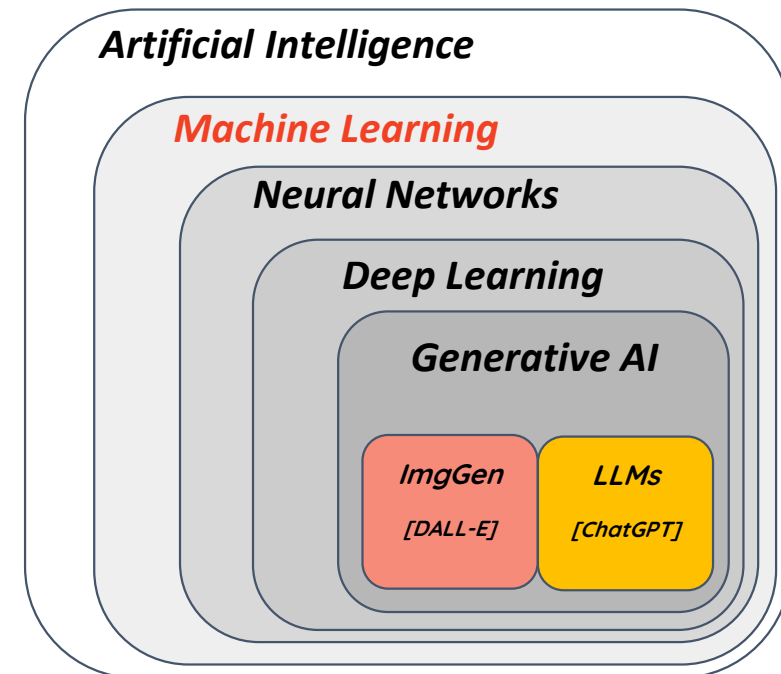
ChatGPT: It is is a type of AI that allows software applications to become more accurate at predicting outcomes without being explicitly programmed. It is based on the idea that systems can learn from data, identify patterns and make decisions with minimal human intervention.

In practice...



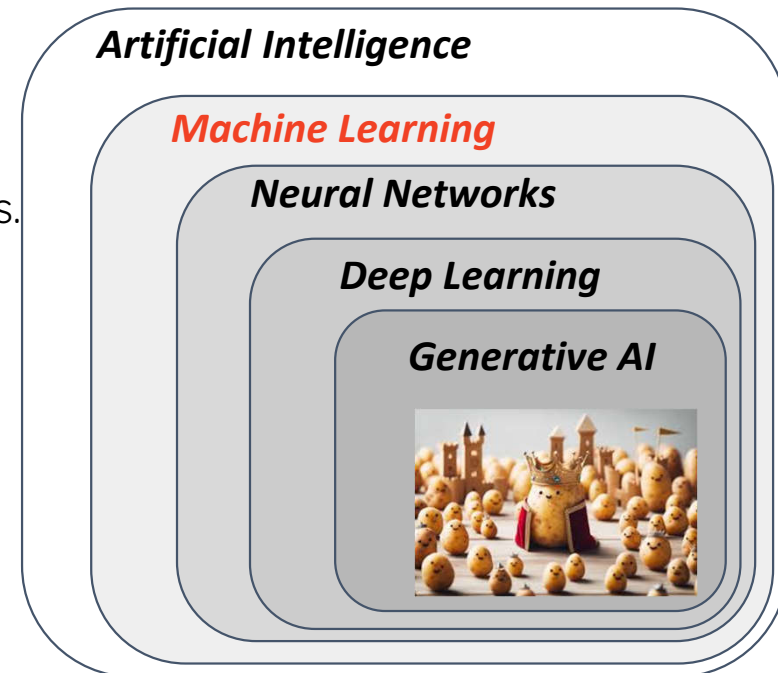
Machine-learning - what, why, and when?

- What is Machine learning (ML)?
 - ML is the study of computer algorithms that improve automatically through experience and by the use of data.



Machine-learning - what, why, and when?

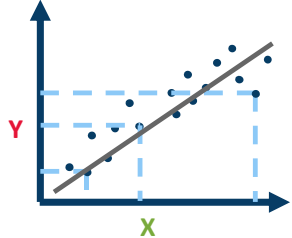
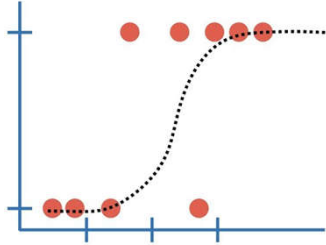
- What is Machine learning (ML)?
 - ML is the study of computer algorithms that improve automatically through experience and by the use of data.
- Why is it useful - especially in life sciences?
 - Biology, Medicine, Environmental sciences comprise phenomena (e.g. a disease) with large number of variables.
 - We want to model complex relationships within these variables and make accurate predictions.



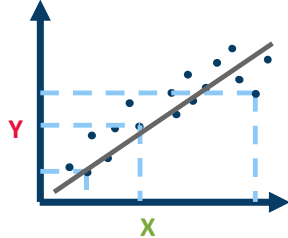
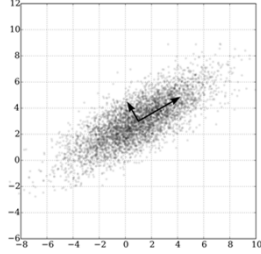
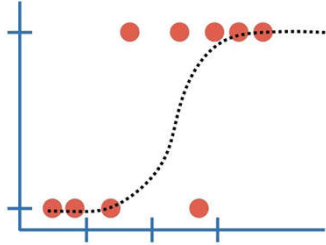
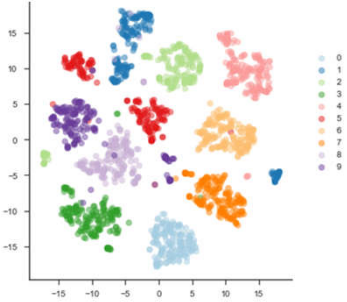
Machine-learning - what, why, and when?

- What is Machine learning (ML)?
 - ML is the study of computer algorithms that improve automatically through experience and by the use of data.
- Why is it useful - especially in life sciences?
 - Biology, Medicine, Environmental sciences comprise phenomena (e.g. a disease) with large number of variables.
 - We want to model complex relationships within these variables.
- When do I use it?
 - You are interested in 1) prediction tasks or 2) low-dimensional representation.
 - You have sufficient data.

Types of ML Algorithms

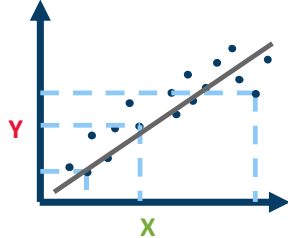
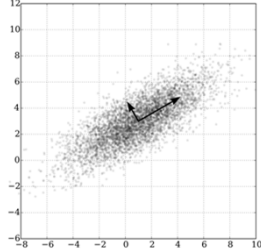
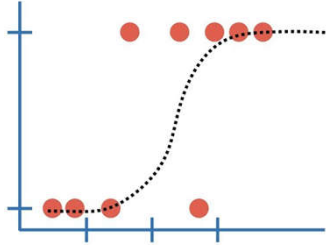
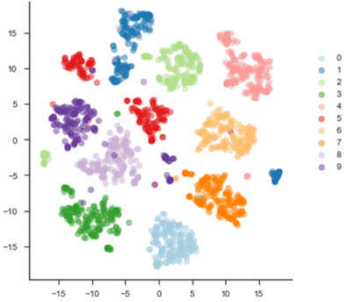
Outcome	Supervised Learning	
Continuous	Regression 	
Categorical	Classification 	

Types of ML Algorithms


Outcome	Supervised Learning	Unsupervised Learning
Continuous	Regression 	Principal Component 
Categorical	Classification 	Clustering 

** ... and Reinforcement Learning*


Types of ML Algorithms

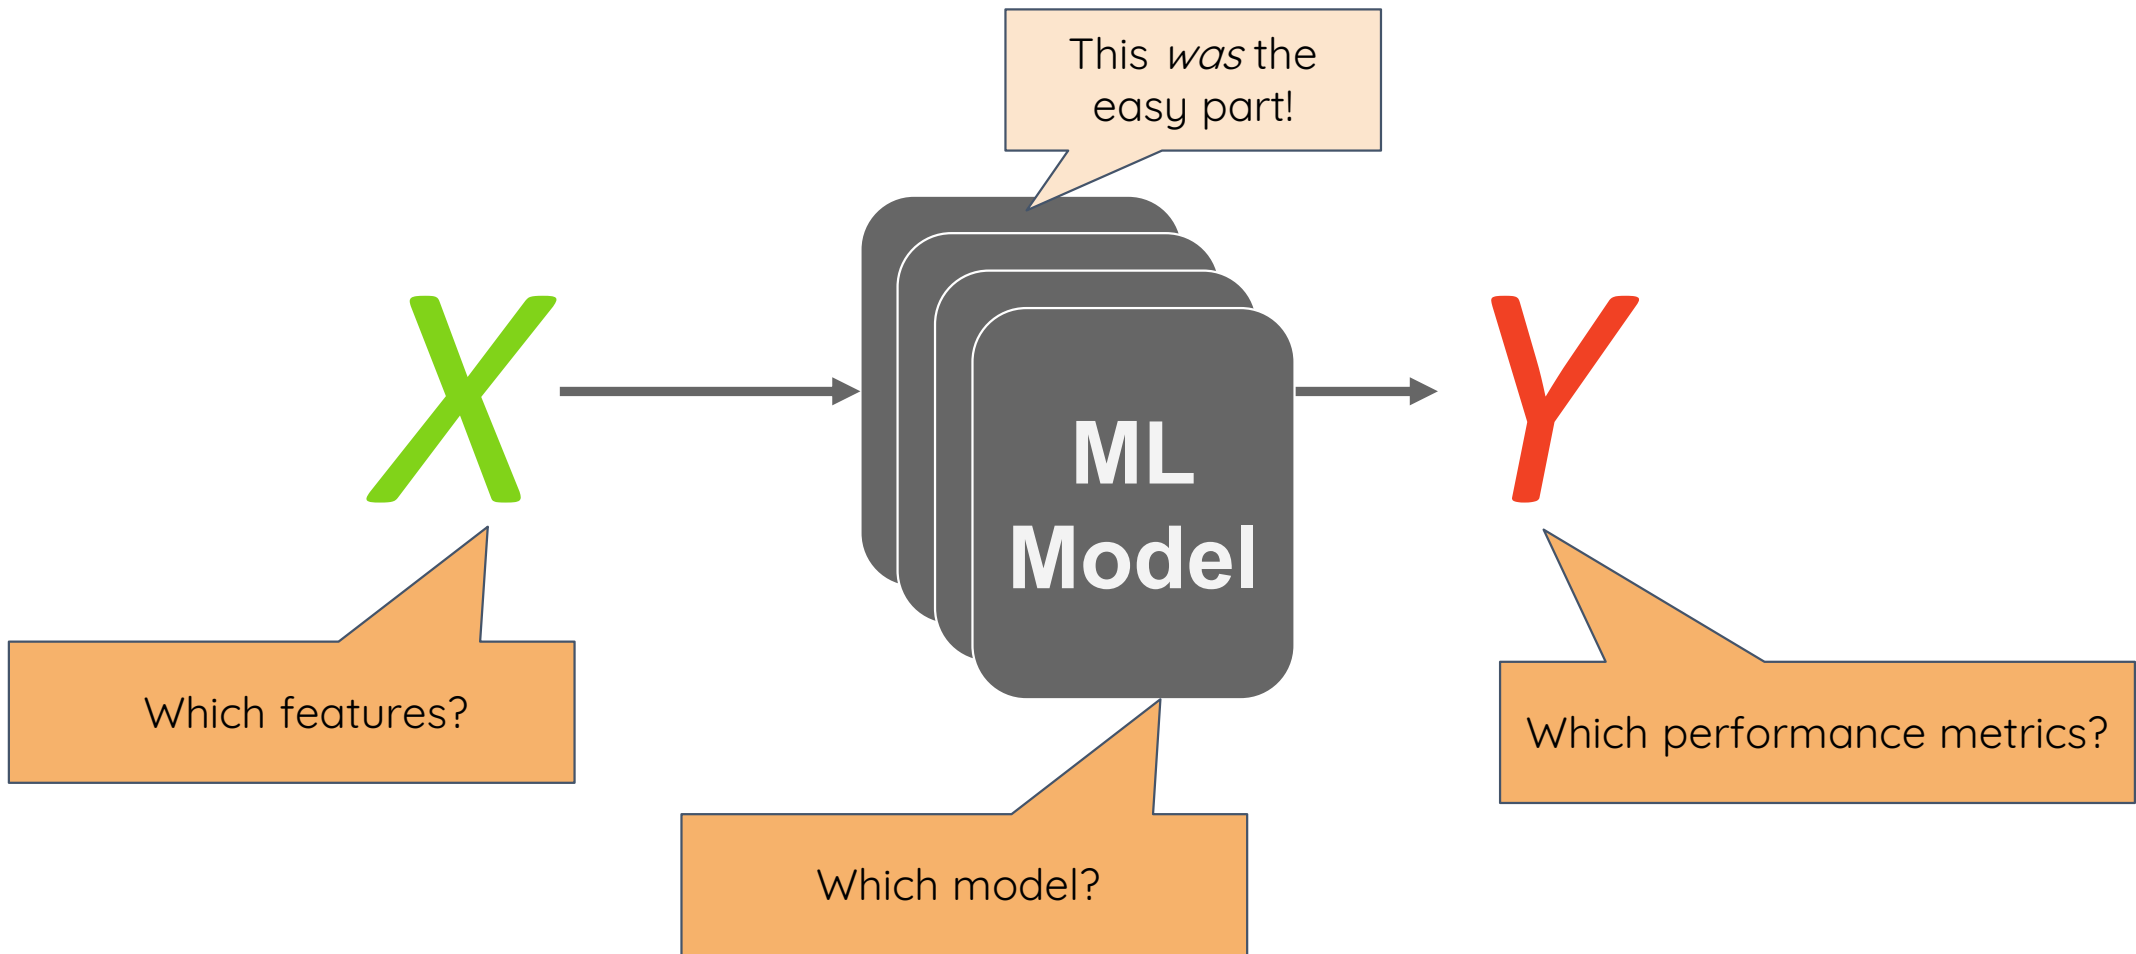
Outcome	Supervised Learning	Unsupervised Learning
Continuous	Regression 	Principal Component 
Categorical	Classification 	Clustering 

Training a machine-learning model



Training a machine-learning model






How do I validate my design choices?

X

Which features?



*This was the
easy part!*

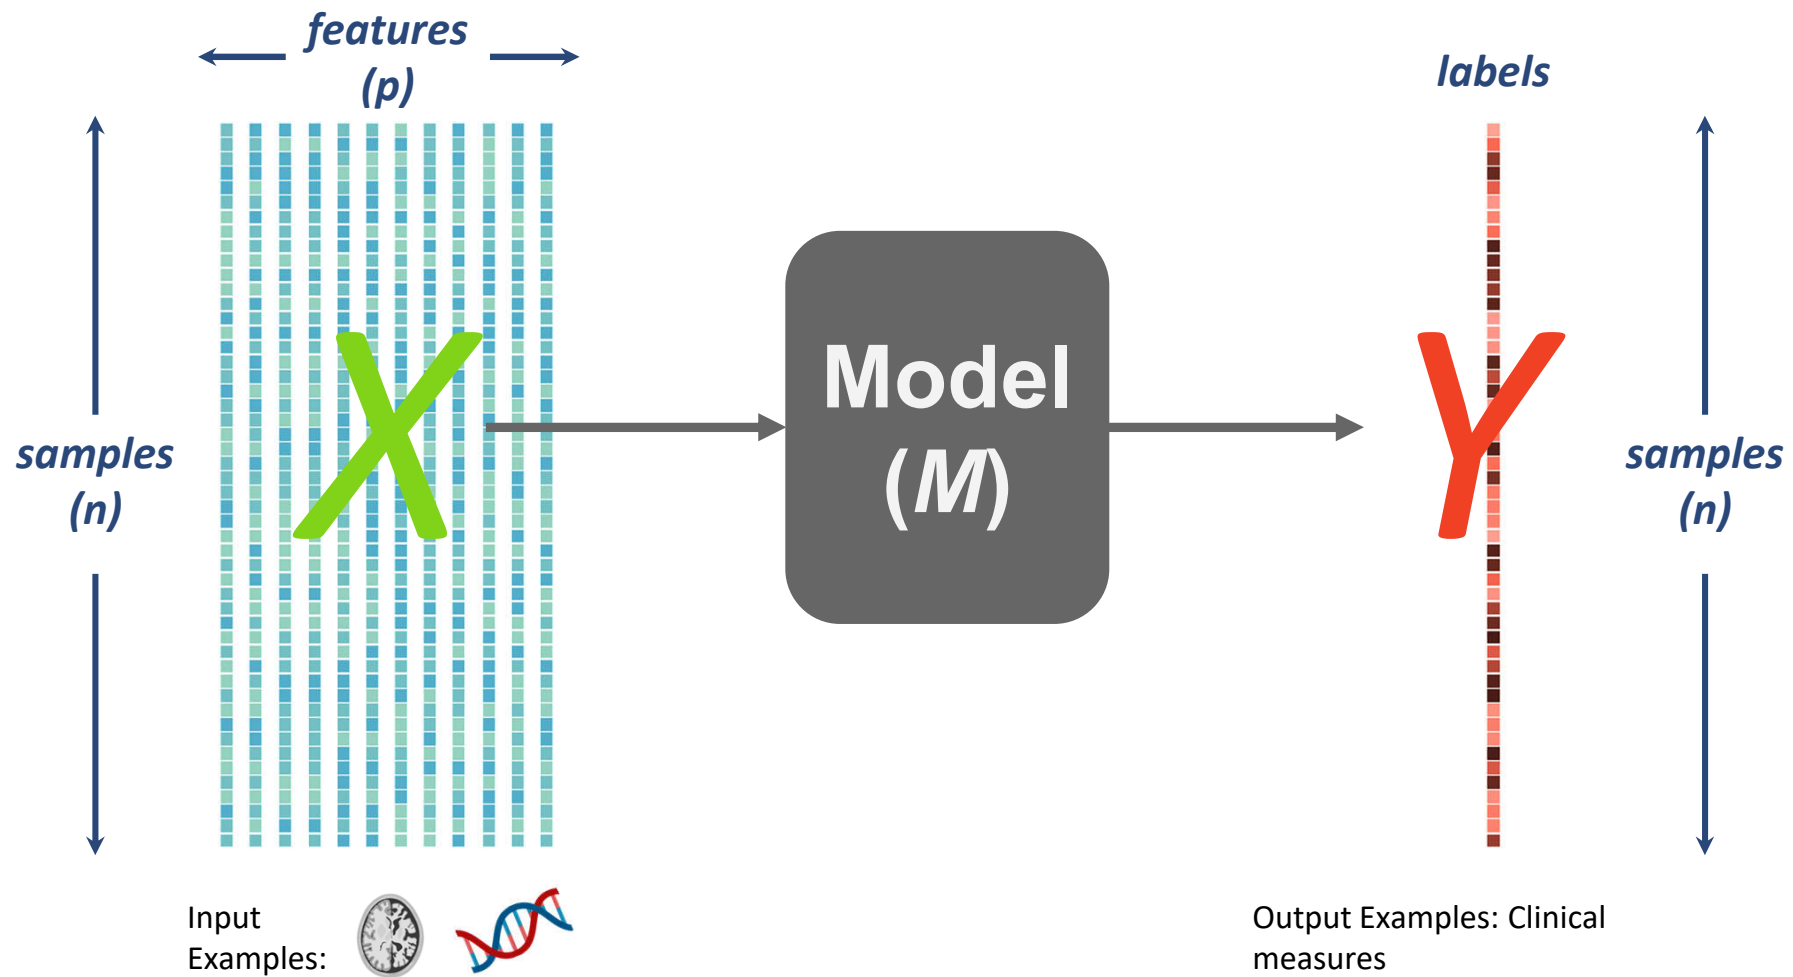


Y

Which performance metrics?

Which model?

Terminology




Supervised Learning: Models

- Goal: *Learn* parameters (or weights) of a model (M) that maps x to y

Supervised Learning: Models


- Goal: *Learn* parameters (or weights) of a model (M) that maps \mathbf{x} to \mathbf{y}
- Example models:
 - Linear / Logistic regression




Linear Regression

Supervised Learning: Models

- Goal: *Learn* parameters (or weights) of a model (M) that maps x to y
- Example models:
 - Linear / Logistic regression
 - Support vector machines




Linear Regression




SVM

Supervised Learning: Models


- Goal: *Learn* parameters (or weights) of a model (M) that maps **x** to **y**
- Example models:
 - Linear / Logistic regression
 - Support vector machines
 - Tree-ensembles: random forests, gradient boosting



Linear Regression




SVM




Supervised Learning: Models


- Goal: *Learn* parameters (or weights) of a model (M) that maps **x** to **y**
- Example models:
 - Linear / Logistic regression
 - Support vector machines
 - Tree-ensembles: random forests, gradient boosting
 - Artificial Neural networks




Linear Regression



SVM





Tree-ensembles



ANN

Model Fitting

- How do we learn the model weights?
 - Example: Linear regression
 - Model: $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2$
 - Loss function: $\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
 - Optimization: Gradient descent





Model Fitting

- Gradient descent with a **single** input variable and **n** samples

- Start with random weights (β_0 and β_1)
- Compute loss (i.e. MSE)
- Update weights based on the gradient


$$\hat{y}_i = \beta_0 + \beta_1 x_i$$

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



Model Fitting

- Gradient descent for complex models with non-convex loss functions
 - Start with random weights (β_0 and β_1)
 - Compute loss
 - Update weights based on the gradient



Model Fitting

- Can we control this fitting process to get a model with specific characteristics?

Model Fitting

- Can we control this fitting process to get a model with specific characteristics?
 - We have strong prior beliefs about what is a **plausible** model
 - e.g. I believe a disease symptom can be predicted with few genes.
 - Practical reasons
 - Prevent overfitting ($n_{\text{features}} \gg n_{\text{samples}}$)

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_{p-1} x_{p-1} + \beta_p x_p$$

Model Fitting

- Can we control this fitting process to get a model with specific characteristics?
 - We have strong prior beliefs about what is a **plausible** model
 - e.g. I believe a disease symptom can be predicted with few genes.
 - Practical reasons
 - **Prevent overfitting ($n_{\text{features}} \gg n_{\text{samples}}$)**
- **Yes! → Model regularization**

Model Fitting: Regularization

- How do we do it?
 - Modify the loss function
 - Constrain the learning process
- Examples:
 - L1 i.e. Lasso
 - L2 i.e. Ridge

Model Fitting: Regularization

- How do we do it?
 - Modify the loss function
 - Constrain the learning process
- Examples:
 - **L1 i.e. Lasso**
 - **L2 i.e. Ridge**

- 1) L1/Lasso: constrains parameters to be *sparse*

$$\text{MSE} = \sum_{i=1}^n (y_i - \underbrace{[\beta_0 + \sum_{j=1}^p x_{ij} \beta_j]}_{\hat{y}_i})^2 + \underbrace{\lambda \sum_{j=1}^p |\beta_j|}_{L_1}$$

- 1) L2/Ridge: constrains parameters to be *small*

$$\text{MSE} = \sum_{i=1}^n (y_i - \underbrace{[\beta_0 + \sum_{j=1}^p x_{ij} \beta_j]}_{\hat{y}_i})^2 + \underbrace{\lambda \sum_{j=1}^p \beta_j^2}_{L_2}$$

Model Fitting: Scikit-learn syntax

```
# import
```

```
from sklearn import linear_model, svm
```

```
# data
```

```
X = [[0, 0], [1, 1]]
```

```
y = [0, 1]
```


Model Fitting: Scikit-learn syntax

```
# import
from sklearn import linear_model, svm

# data
X = [[0, 0], [1, 1]]
y = [0, 1]

# pick a model
model = linear_model.Lasso(alpha=0.1) # model = svm.SVC()

# fit the model with data
model.fit(X, y)
```

Model Fitting: Scikit-learn syntax

```
# import
from sklearn import linear_model, svm

# data
X = [[0, 0], [1, 1]]
y = [0, 1]


# pick a model
model = linear_model.Lasso(alpha=0.1) # model = svm.SVC()

# fit the model with data
model.fit(X, y)

# predict on new data
y_pred = model.predict([[1, 0]])
```


Model Evaluation

- Is the model generalizable?
- How do we sample train and test sets?
- How do we select a model?





Model Evaluation

- **Is the model generalizable?**
- How do we sample train and test sets?
- How do we select a model?





Model Evaluation

- Train performance \neq Test performance
 - Model: Underfitting vs Overfitting
 - Errors: Bias - Variance tradeoff




Model Evaluation


- Train performance \neq Test performance
 - Model: Underfitting vs Overfitting
 - Errors: Bias - Variance tradeoff
 - Regression example



Underfitting



Optimal





Overfitting


● Train set

Model Evaluation


- Train performance \neq Test performance
 - Model: Underfitting vs Overfitting
 - Errors: Bias - Variance tradeoff
 - Regression example



Underfitting



Optimal





Overfitting

● Train set


● Test set

Model Evaluation


- Train performance \neq Test performance
 - Model: Underfitting vs Overfitting
 - Errors: Bias - Variance tradeoff
 - Classification example (i.e. separate “o” vs “x”)



Underfitting



Optimal





Overfitting

● Train class_1


✕ Train class_2

Model Evaluation


- Train performance \neq Test performance
 - Model: Underfitting vs Overfitting
 - Errors: Bias - Variance tradeoff
 - Classification example (i.e. separate “o” vs “x”)



Underfitting



Optimal



Overfitting

○ Train class_1

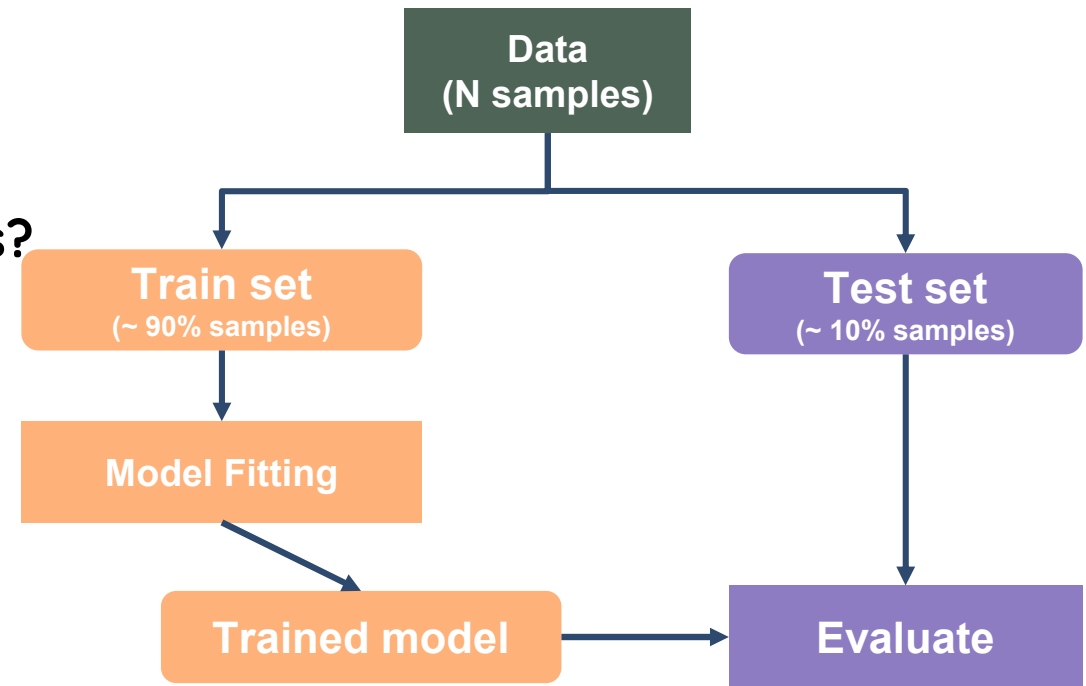
✕ Train class_2

● Test class_1

✕ Test class_2


Model Evaluation

- Is the model generalizable?
- **How do we sample train and test sets?**
- How do we select a model?



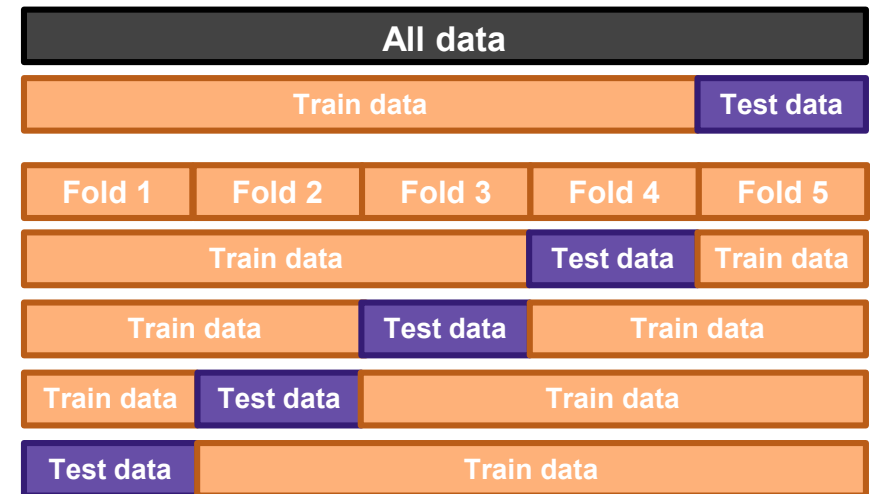
Model Evaluation: Cross-Validation (Outer loop)

- How do we sample train and test sets?
 - Train set: learn model parameters
 - Test set (a.k.a held-out sample): Evaluate model performance



Model Evaluation: Cross-Validation (Outer loop)

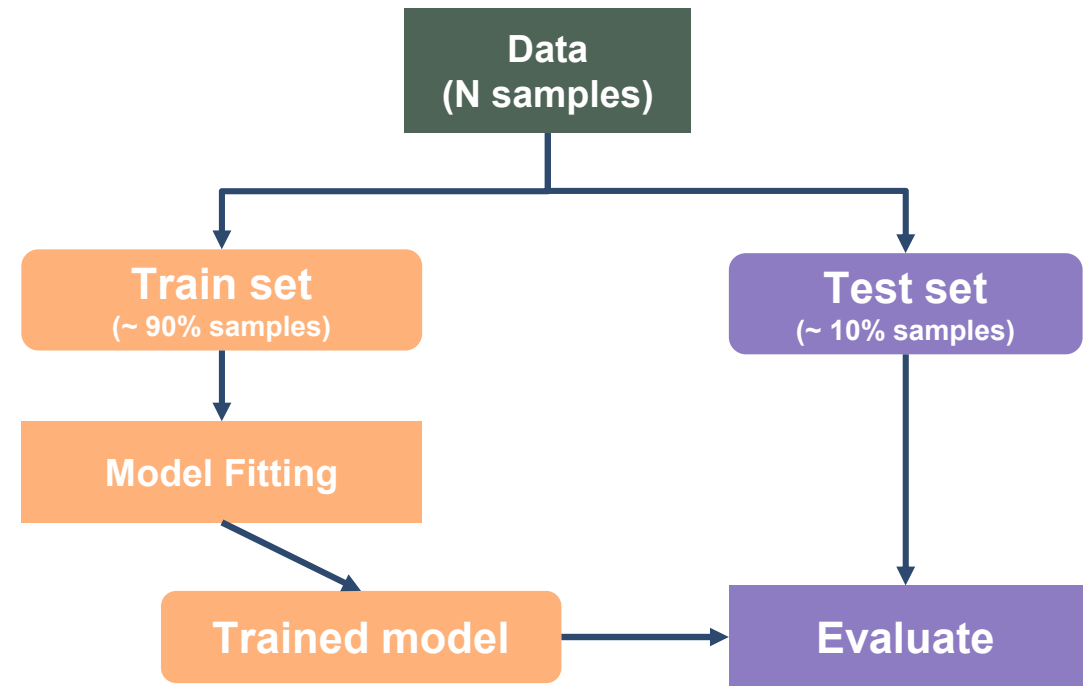
- How do we sample train and test sets?
 - Train set: learn model parameters
 - Test set (a.k.a held-out sample): Evaluate model performance
 - Repeat for different Train-Test splits
 - k-fold, shuffle-split
 - Report performance statistics over all test folds



CV outer loop

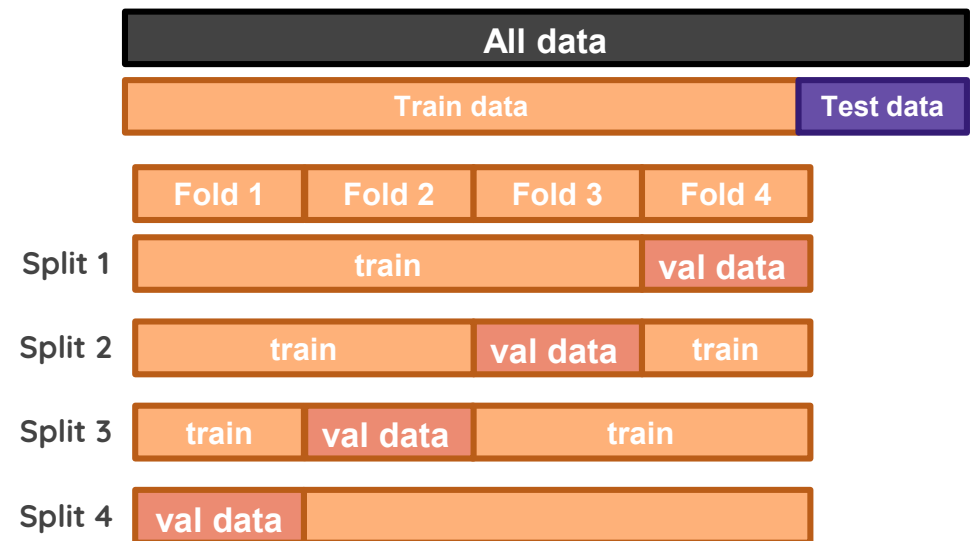
Model Evaluation

- Is the model generalizable?
- How do we sample train and test sets?
- **How do we select a model?**



Model Evaluation: Cross-Validation (Inner loop)

- How do we select a model?
 - Tune *hyper-parameters* of a model
 - Compare several different model architectures
 - Select / transform raw features
- This repeats for all train-test splits in the outer loop



CV inner loop

Model Evaluation: Hyper-parameters

- Hyper-parameter \neq parameter (or weights)
 - Parameters are **learned**; hyper-parameters are **chosen**!

Model Evaluation: Hyper-parameters

- Hyper-parameter \neq parameter (or weights)
 - Parameters are **learned**; hyper-parameters are **chosen**!
- Examples:
 - Degree of model (eg. linear vs quadratic)
 - Kernels
 - Number of trees
 - Number of layers, filters, batch-size, learning-rate in ANNs

Model Evaluation: Hyper-parameters


- Hyper-parameter \neq parameter (or weights)
 - Parameters are **learned**; hyper-parameters are **chosen**!
- Examples:
 - Degree of model (eg. linear vs quadratic)
 - Kernels
 - Number of trees
 - Number of layers, filters, batch-size, learning-rate in ANNs
- How do we choose them?
 - Prior beliefs \rightarrow eg. cortical thickness and age have quadratic relationship.
 - Arbitrarily \rightarrow we gotta start with something!
 - Trial and error \rightarrow do a computationally feasible grid-search.

Performance Scores


- Loss functions → computationally well-suited metrics
 - May / need not completely capture performance metrics of interest
- Scores → practically useful metrics
 - Binary classification

Confusion Matrix		Ground Truth	
		POSITIVE	NEGATIVE
Prediction	POSITIVE	TP	FP
	NEGATIVE	FN	TN

False Positive



False Negative



Performance Scores

- ML model that detects Covid from chest CTs. Current Covid prevalence ~ 1%.
 - FP: model predicts *Covid* when person is *healthy*
 - FN: model predicts *healthy* when person has *Covid*
- What happens if we build model that predicts everyone as healthy?
 - i.e. zero FPs!



Performance Scores

- ML model that detects Covid from chest CTs. Current Covid prevalence ~ 1%.
 - FP: model predicts *Covid* when person is *healthy*
 - FN: model predicts *healthy* when person has *Covid*
- What happens if we build model that predicts everyone as healthy?

Score	Formula	Null	What does it tell us?	When do I use it?
Accuracy	$(TP+TN) / (TP+FP+FN+TN)$	0.99	How many people did we correctly predict out of all the people scanned?	FNs & FPs have similar costs
Precision (i.e. PPV)	$TP/(TP+FP)$	NaN	How many of those who we predicted as “covid” do actually have “covid”?	If you want to be more confident of your TPs
Recall (aka Sensitivity)	$TP/(TP+FN)$	0	Of all the people who have covid, how many of those did we correctly predict?	If you prefer FPs over FNs.
Specificity	$TN/(TN+FP)$	1	Of all the people who are healthy, how many of those did we correctly predict?	If you prefer FNs over FPs.
F1	$2*(Recall * Precision) / (Recall + Precision)$	NaN	Harmonic mean(average) of the precision and recall.	When you have an uneven class distribution

Performance Curves

- Receiver Operating Characteristic (ROC) → Want high area-under-the-curve (AUC)
- Precision-Recall → Want high AUC or high Average precision (AP)



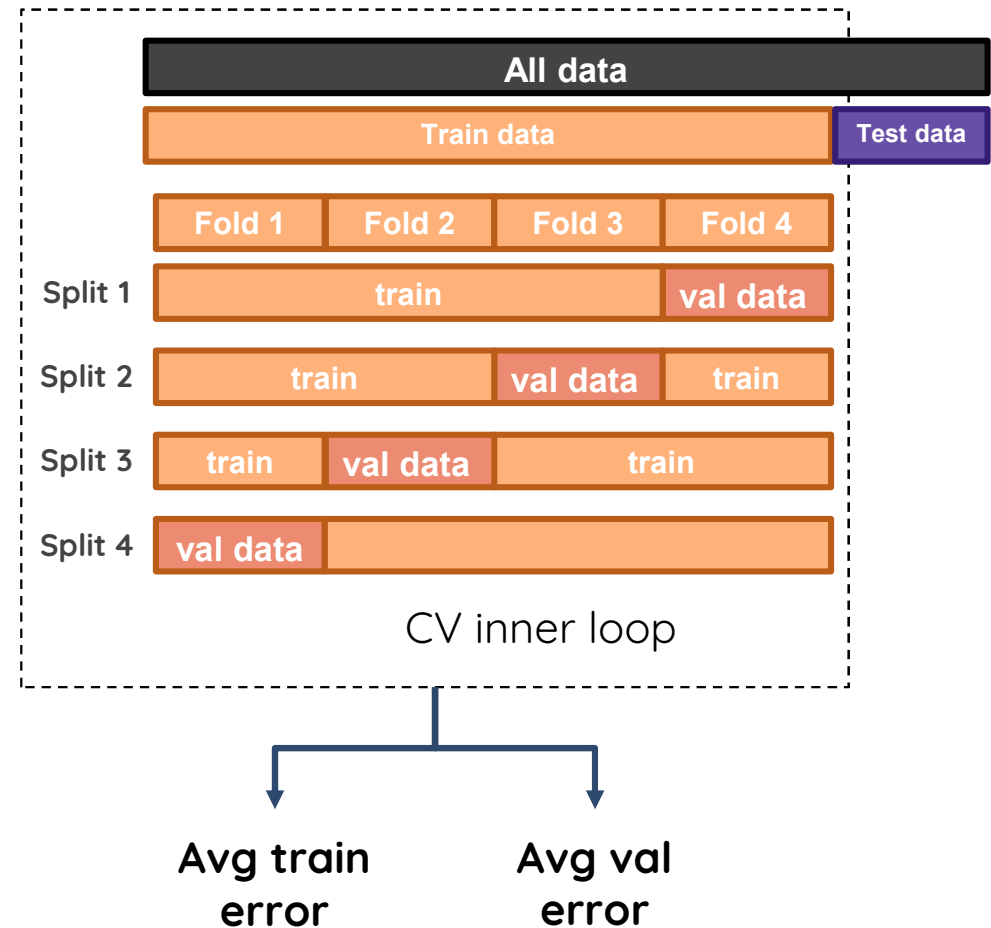
Practical intuition

Task: Segmentation, diagnosis etc

- Human error ~ 2%
 - Train error ~ 10%
 - Val error ~ 20%
- Bias / underfit
- Variance / overfit

What do we do?

- Underfitting → Bigger/different model
- Overfitting → More data / regularization



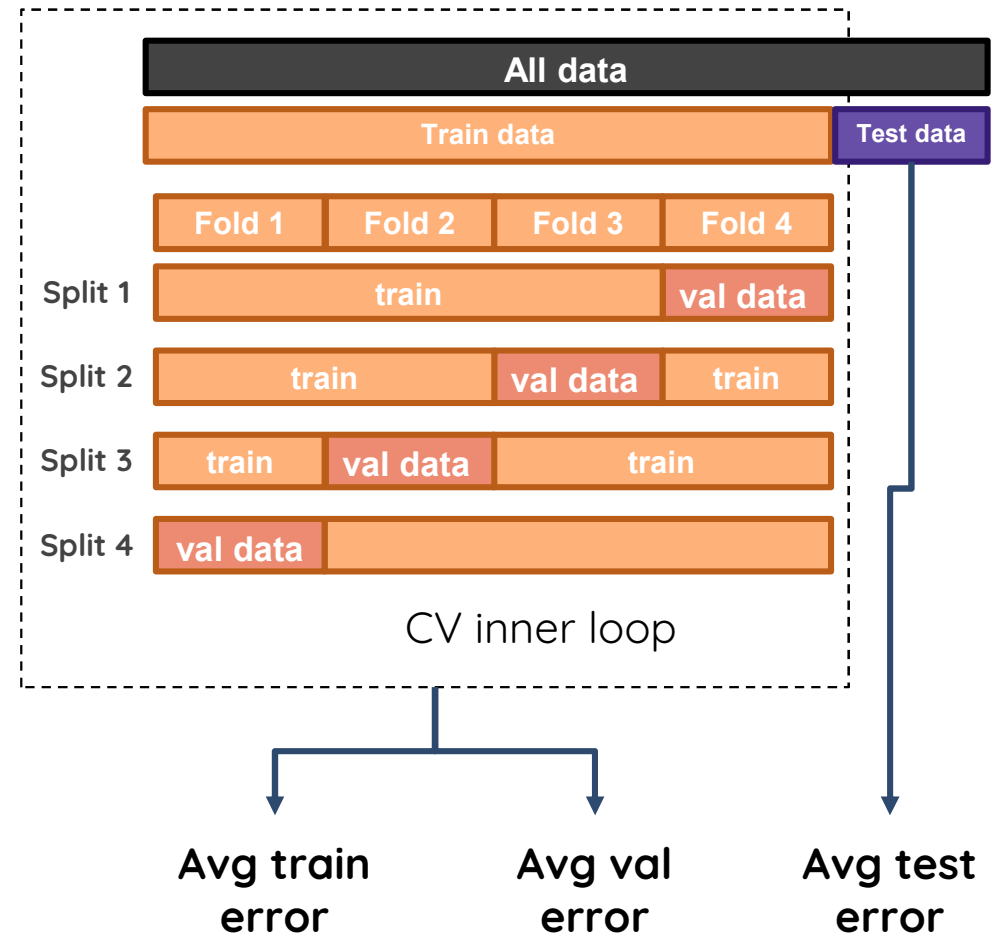
Practical intuition

Task: Segmentation, diagnosis etc

- Human error ~ 2%
 - Train error ~ 5%
 - Val error ~ 5%
 - Test error ~ 20%
- } dataset shift (overfit)


What do we do?

- Feature shift → get / generate more data
- Concept drift → fix / refine labels



Deep-learning

- Why the buzz?
 - Works amazing on spatio-temporal input
 - Highly flexible → universal function approximator



ANN for handwritten-digit images
(gif source: [3b1b](#))

Deep-learning


- Why the buzz?
 - Works amazing on spatio-temporal input
 - Highly flexible → universal function approximator
- What are the challenges?
 - Large number of parameters (175B!) → data hungry
 - Large number of hyper-parameters → difficult to train



LLM Transformers
(gif source: [3b1b](#))

Deep-learning


- Why the buzz?
 - Works amazing on spatio-temporal input
 - Highly flexible → universal function approximator
- What are the challenges?
 - Large number of parameters (175B!) → data hungry
 - Large number of hyper-parameters → difficult to train
- When do I use it?
 - If you have highly-structured input, eg. medical images.
 - You have a lot of data and computational resources.



Source:
<https://github.com/fepegar/torchio>

Pitfalls and Challenges


- Models do not generalize even after good CV performance
 - Implicit double-dipping
 - Dataset biases (eg. North-American demographics)
 - Noisy labels (eg. diagnosis definitions)
 - Data distribution shifts (eg. assay, scanner upgrades)



Pitfalls and Challenges

- Models do not generalize even after good CV performance

- Implicit double-dipping
- Dataset biases (eg. North-American demographics)
- Noisy labels (eg. diagnosis definitions)
- Data distribution shifts (eg. assay, scanner upgrades)



- Unnecessary complexity

- Do I really need a giant deep-net or a simple linear model would do?



ML Novice Checklist

- Data

- What is my n_features and n_samples?
- Am I [encoding](#) categorical data correctly?
- Am I using information (e.g. mean) from test set to preprocess (eg. zscore) the data?

ML Novice Checklist

○ Data


- What is my n_features and n_samples?
- Am I [encoding](#) categorical data correctly?
- Am I using information (e.g. mean) from test set to preprocess (eg. zscore) the data?

○ Model

- Do my performance metrics capture the practical use-case of interest?
- What is the null / dummy model performance?
 - Classification: Predict majority class all the time
 - Regression: Predict the median value all the time
- Am I interpreting model parameters (i.e. weights) correctly?

Takeaways

- Supervised ML is useful for **predictions** but **not really for explanations**
 - eg. image segmentation, prognosis, drug development
- Our job is to ensure **generalizability** of these models
 - Multitude of validations
 - Understanding model biases and limitations
- ***Engineering tools*** vs *Scientific discovery*
 - Interpretability and explainability



Explainable AI

Food for thought

- Ethical dilemmas
- Societal implications
- What's real?



Ceci n'est pas un pape

ML for Neuroimaging

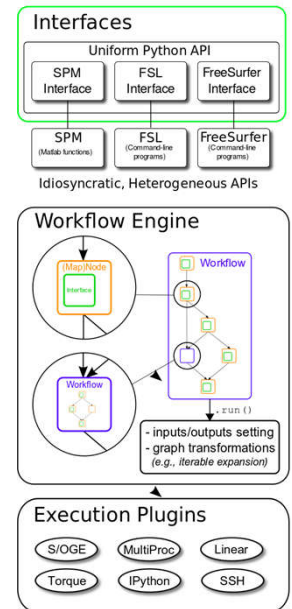
- Why Neuroimaging is special?
 - Imaging data is huge (Dimension reduction)
 - Imaging data has multi-modality (Fusion)
 - Imaging data is noisy (indirect measure of brain activities)

Python tools for Neuroimaging ML

- **Community of practice:** Nipy.org
- Pipelines and interfaces: Nipype
- Anatomy: dipy, Mindboggle
- File I/O and data management: nibabel, Scitran SDM, pybids...
- fMRI: Nipy, Nitime, popeye...
- ML: **nilearn**, PyMVPA...
- i/S/M/EEG: MNE...
- Visualization: napari-nibabel, niwidgets...



Nipype: Neuroimaging in Python Pipelines and Interfaces




<https://nipype.readthedocs.io/en/latest/index.html>
<https://nilearn.github.io/dev/index.html>

Nilearn



Nilearn enables approachable and versatile analyses of brain volumes. It provides **statistical and machine-learning tools**, with instructive documentation & open community.

It supports general linear model (GLM) based analysis and leverages the scikit-learn Python toolbox for multivariate statistics with applications such as predictive modelling, classification, decoding, or connectivity analysis.



Nilearn example: Extracting features from Imaging data

```
from nilearn import datasets

development_dataset = datasets.fetch_development_fmri(n_subjects=30)
```

```
import nibabel as nib

# Subset to just the first image
img = nib.load(development_dataset.func[0])
img.shape
```

```
import numpy as np

msdl_atlas = datasets.fetch_atlas_msdl()

msdl_coords = msdl_atlas.region_coords
n_regions = len(msdl_coords)

print(f'MSDL has {n_regions} ROIs, part of the following networks :\n{np.unique(msdl_atlas.network
```

```
from nilearn import input_data

masker = input_data.NiftiMapsMasker(
    msdl_atlas.maps, resampling_target="data",
    t_r=2, detrend=True,
    low_pass=0.1, high_pass=0.01).fit()
```

```
roi_time_series = masker.transform(development_dataset.func[0])
roi_time_series.shape
```

```
import pandas as pd

pd.read_table(development_dataset.confounds[0]).head()
```

nilearn.maskers: Extracting Signals from Brain Images

The `nilearn.maskers` contains masker objects.


User guide: See the `NiftiMasker` applying a mask to load time-series section for further details.

Classes:

<code>BaseMasker</code> (0)	Base class for NiftiMaskers.
<code>NiftiMasker</code> (mask_img, runs, ...)	Applying a mask to extract time-series from Nimg-like objects.
<code>NiftiSpheresMasker</code> (mask_img, smoothing_fwhm, ...)	Applying a mask to extract time-series from multiple Nimg-like objects.
<code>NiftiLabelsMasker</code> (labels_img, labels, ...)	Class for extracting data from Nimg-like objects using labels of non-overlapping brain regions.
<code>NiftiMapsMasker</code> (labels_img, labels, ...)	Class for extracting data from multiple Nimg-like objects using labels of non-overlapping brain regions.
<code>NiftiMapsMasker</code> (maps_img, mask_img, ...)	Class for extracting data from Nimg-like objects using maps of potentially overlapping brain regions.
<code>NiftiMapsMasker</code> (maps_img, mask_img, ...)	Class for extracting data from multiple Nimg-like objects using maps of potentially overlapping brain regions.
<code>NiftiSpheresMasker</code> (seeds, radius, ...)	Class for masking of Nimg-like objects using seeds.

• Examples/masker_reports

- Nifti masker
- Nifti labels masker
- Nifti maps masker
- Nifti sphere masker



```
corrected_roi_time_series = masker.transform(
    development_dataset.func[0], confounds=development_dataset.confounds[0])
corrected_correlation_matrix = correlation_measure.fit_transform(
    [corrected_roi_time_series])[0]
np.fill_diagonal(corrected_correlation_matrix, 0)
plotting.plot_matrix(corrected_correlation_matrix, labels=msdl_atlas.labels,
    vmax=0.8, vmin=-0.8, colorbar=True)
```

Nilearn example: functional connectivity



```
import numpy as np
import matplotlib.pyplot as plt
from nilearn import (datasets, input_data, plotting)
from nilearn.connectome import ConnectivityMeasure

development_dataset = datasets.fetch_development_fmri(n_subjects=30)
msdl_atlas = datasets.fetch_atlas_msdl()

masker = input_data.NiftiMapsMasker(
    msdl_atlas.maps, resampling_target="data",
    t_r=2, detrend=True,
    low_pass=0.1, high_pass=0.01).fit()
correlation_measure = ConnectivityMeasure(kind='correlation')
```

```
children = []
pooled_subjects = []
groups = [] # child or adult

for func_file, confound_file, phenotypic in zip(
    development_dataset.func,
    development_dataset.confounds,
    development_dataset.phenotypic):

    time_series = masker.transform(func_file, confounds=confound_file)
    pooled_subjects.append(time_series)

    if phenotypic['Child_Adult'] == 'child':
        children.append(time_series)


    groups.append(phenotypic['Child_Adult'])

print('Data has {} children.'.format(len(children)))
```




```
from sklearn.metrics import accuracy_score
from sklearn.model_selection import StratifiedShuffleSplit
from sklearn.svm import LinearSVC

kinds = ['correlation', 'partial correlation', 'tangent']
_, classes = np.unique(groups, return_inverse=True)
cv = StratifiedShuffleSplit(n_splits=15, random_state=0, test_size=5)
pooled_subjects = np.asarray(pooled_subjects)
```

Control Covariance



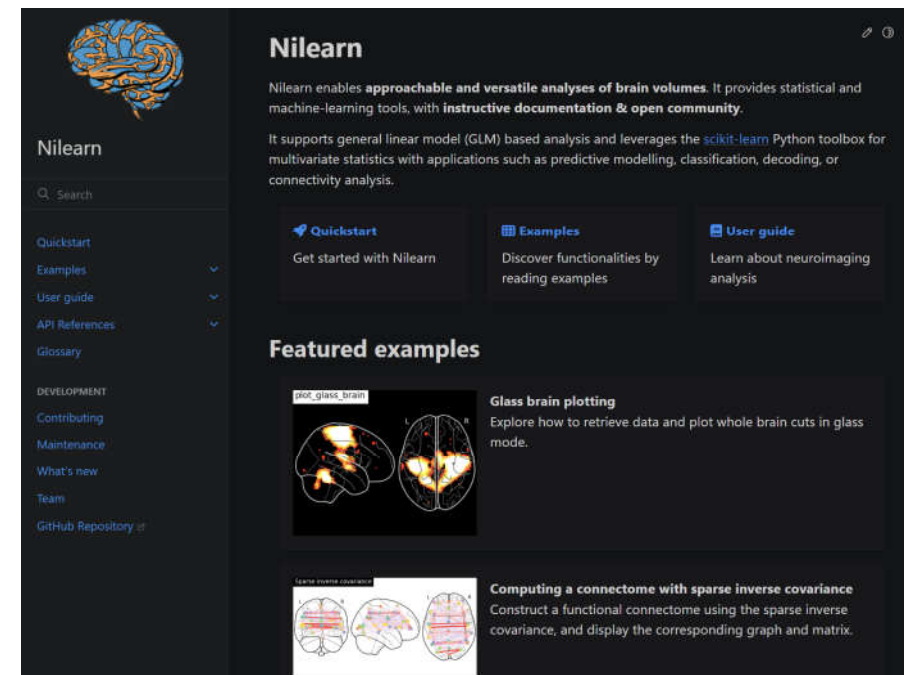
Control Sparse inverse covariance (GraphicalLasso)



```
scores = {}
for kind in kinds:
    scores[kind] = []
    for train, test in cv.split(pooled_subjects, classes):
        # "ConnectivityMeasure" can output the estimated subjects' coefficients
        # as a 1D arrays through the parameter "vectorize".
        connectivity = ConnectivityMeasure(kind=kind, vectorize=True)
        # build vectorized connectomes for subjects in the train set
        connectomes = connectivity.fit_transform(pooled_subjects[train])
        # fit the classifier
        classifier = LinearSVC().fit(connectomes, classes[train])
        # make predictions for the left-out test subjects
        predictions = classifier.predict(
            connectivity.transform(pooled_subjects[test]))
        # store the accuracy for this cross-validation fold
        scores[kind].append(accuracy_score(classes[test], predictions))
```

Takeaways

- Python based env are ready for neuroimaging studies
- Learning by doing and start with the examples
- ***Engineering tools*** vs *Scientific discovery*
 - Interpretability and explainability



Useful resources

McGill QLS612 course: <https://neurodatascience.github.io/QLS612-Overview/>

https://inria.github.io/scikit-learn-mooc/ml_concepts/slides.html

<https://www.3blue1brown.com/topics/linear-algebra>

3b1b Gradient Descent: <https://www.youtube.com/watch?v=IHZwWFHWa-w>

Python Neuroimaging libs family: <https://nipy.org/>